

```
In [1]: #import libraries

import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
%matplotlib inline
```

## PROJECT DESCRIPTION:

Heart failure occurs when the heart muscle doesn't pump blood well as it should. Cardiovascular diseases (CVDs) are the number one cause of death globally, taking an estimated 17.9 million lives each year, which account to 31% of all deaths worldwide.

This project is focused on the 12 features that can be used to predict mortality by heart failure. Data was collected on these 12 features for the analysis.

## Description of attributes in failure dataset

- age: Age of the patient
- anaemia: Haemoglobin level of patient (Boolean)
- creatinine\_phosphokinase: Level of the CPK enzyme in the blood (mcg/L)
- diabetes: If the patient has diabetes (Boolean)
- ejection\_fraction: Percentage of blood leaving the heart at each contraction
- high\_blood\_pressure: If the patient has hypertension (Boolean)
- platelets: Platelet count of blood (kiloplatelets/mL)
- serum\_creatinine: Level of serum creatinine in the blood (mg/dL)
- serum\_sodium: Level of serum sodium in the blood (mEq/L)
- sex: Sex of the patient
- smoking: If the patient smokes or not (Boolean)
- time: Follow-up period (days)
- DEATH\_EVENT: If the patient deceased during the follow-up period (Boolean)
- For attributes with boolean values; 0 means negative while 1 means positive

In [2]: *#let's read our data set*

```
heart_failure = pd.read_csv('heart_failure_clinical_records_dataset.csv')
heart_failure
```

Out[2]:

|     | age  | anaemia | creatinine_phosphokinase | diabetes | ejection_fraction | high_blood_pressure | platelets | serum_creatinine | serum_sodium | sex | si |
|-----|------|---------|--------------------------|----------|-------------------|---------------------|-----------|------------------|--------------|-----|----|
| 0   | 75.0 | 0       | 582                      | 0        | 20                | 1                   | 265000.00 | 1.9              | 130          | 1   |    |
| 1   | 55.0 | 0       | 7861                     | 0        | 38                | 0                   | 263358.03 | 1.1              | 136          | 1   |    |
| 2   | 65.0 | 0       | 146                      | 0        | 20                | 0                   | 162000.00 | 1.3              | 129          | 1   |    |
| 3   | 50.0 | 1       | 111                      | 0        | 20                | 0                   | 210000.00 | 1.9              | 137          | 1   |    |
| 4   | 65.0 | 1       | 160                      | 1        | 20                | 0                   | 327000.00 | 2.7              | 116          | 0   |    |
| ... | ...  | ...     | ...                      | ...      | ...               | ...                 | ...       | ...              | ...          | ... |    |
| 294 | 62.0 | 0       | 61                       | 1        | 38                | 1                   | 155000.00 | 1.1              | 143          | 1   |    |
| 295 | 55.0 | 0       | 1820                     | 0        | 38                | 0                   | 270000.00 | 1.2              | 139          | 0   |    |
| 296 | 45.0 | 0       | 2060                     | 1        | 60                | 0                   | 742000.00 | 0.8              | 138          | 0   |    |
| 297 | 45.0 | 0       | 2413                     | 0        | 38                | 0                   | 140000.00 | 1.4              | 140          | 1   |    |
| 298 | 50.0 | 0       | 196                      | 0        | 45                | 0                   | 395000.00 | 1.6              | 136          | 1   |    |

299 rows × 13 columns



In [3]: *#check information about our data set*

```
heart_failure.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 299 entries, 0 to 298
```

```
Data columns (total 13 columns):
```

| #  | Column                   | Non-Null Count | Dtype   |
|----|--------------------------|----------------|---------|
| 0  | age                      | 299 non-null   | float64 |
| 1  | anaemia                  | 299 non-null   | int64   |
| 2  | creatinine_phosphokinase | 299 non-null   | int64   |
| 3  | diabetes                 | 299 non-null   | int64   |
| 4  | ejection_fraction        | 299 non-null   | int64   |
| 5  | high_blood_pressure      | 299 non-null   | int64   |
| 6  | platelets                | 299 non-null   | float64 |
| 7  | serum_creatinine         | 299 non-null   | float64 |
| 8  | serum_sodium             | 299 non-null   | int64   |
| 9  | sex                      | 299 non-null   | int64   |
| 10 | smoking                  | 299 non-null   | int64   |
| 11 | time                     | 299 non-null   | int64   |
| 12 | DEATH_EVENT              | 299 non-null   | int64   |

```
dtypes: float64(3), int64(10)
```

```
memory usage: 30.5 KB
```

In [49]: *#checking the summary of our data, let's transpose for quick graphs of our summary*

```
heart_failure.describe().T
```

Out[49]:

|                                 | count | mean          | std          | min     | 25%      | 50%      | 75%      | max      |
|---------------------------------|-------|---------------|--------------|---------|----------|----------|----------|----------|
| <b>age</b>                      | 299.0 | 60.833893     | 11.894809    | 40.0    | 51.0     | 60.0     | 70.0     | 95.0     |
| <b>anaemia</b>                  | 299.0 | 0.431438      | 0.496107     | 0.0     | 0.0      | 0.0      | 1.0      | 1.0      |
| <b>creatinine_phosphokinase</b> | 299.0 | 581.839465    | 970.287881   | 23.0    | 116.5    | 250.0    | 582.0    | 7861.0   |
| <b>diabetes</b>                 | 299.0 | 0.418060      | 0.494067     | 0.0     | 0.0      | 0.0      | 1.0      | 1.0      |
| <b>ejection_fraction</b>        | 299.0 | 38.083612     | 11.834841    | 14.0    | 30.0     | 38.0     | 45.0     | 80.0     |
| <b>high_blood_pressure</b>      | 299.0 | 0.351171      | 0.478136     | 0.0     | 0.0      | 0.0      | 1.0      | 1.0      |
| <b>platelets</b>                | 299.0 | 263358.029264 | 97804.236869 | 25100.0 | 212500.0 | 262000.0 | 303500.0 | 850000.0 |
| <b>serum_creatinine</b>         | 299.0 | 1.393880      | 1.034510     | 0.5     | 0.9      | 1.1      | 1.4      | 9.4      |
| <b>serum_sodium</b>             | 299.0 | 136.625418    | 4.412477     | 113.0   | 134.0    | 137.0    | 140.0    | 148.0    |
| <b>sex</b>                      | 299.0 | 0.648829      | 0.478136     | 0.0     | 0.0      | 1.0      | 1.0      | 1.0      |
| <b>smoking</b>                  | 299.0 | 0.321070      | 0.467670     | 0.0     | 0.0      | 0.0      | 1.0      | 1.0      |
| <b>time</b>                     | 299.0 | 130.260870    | 77.614208    | 4.0     | 73.0     | 115.0    | 203.0    | 285.0    |
| <b>DEATH_EVENT</b>              | 299.0 | 0.321070      | 0.467670     | 0.0     | 0.0      | 0.0      | 1.0      | 1.0      |

In [5]: *#check for the sum of null values*

```
heart_failure.isnull().sum()
```

Out[5]:

|                          |   |
|--------------------------|---|
| age                      | 0 |
| anaemia                  | 0 |
| creatinine_phosphokinase | 0 |
| diabetes                 | 0 |
| ejection_fraction        | 0 |
| high_blood_pressure      | 0 |
| platelets                | 0 |
| serum_creatinine         | 0 |
| serum_sodium             | 0 |
| sex                      | 0 |
| smoking                  | 0 |
| time                     | 0 |
| DEATH_EVENT              | 0 |

dtype: int64

In [6]: *#check the size of the data*

```
heart_failure.shape
```

Out[6]: (299, 13)

In [7]: *#To check for the unique columns in the data set*

```
heart_failure.columns.unique
```

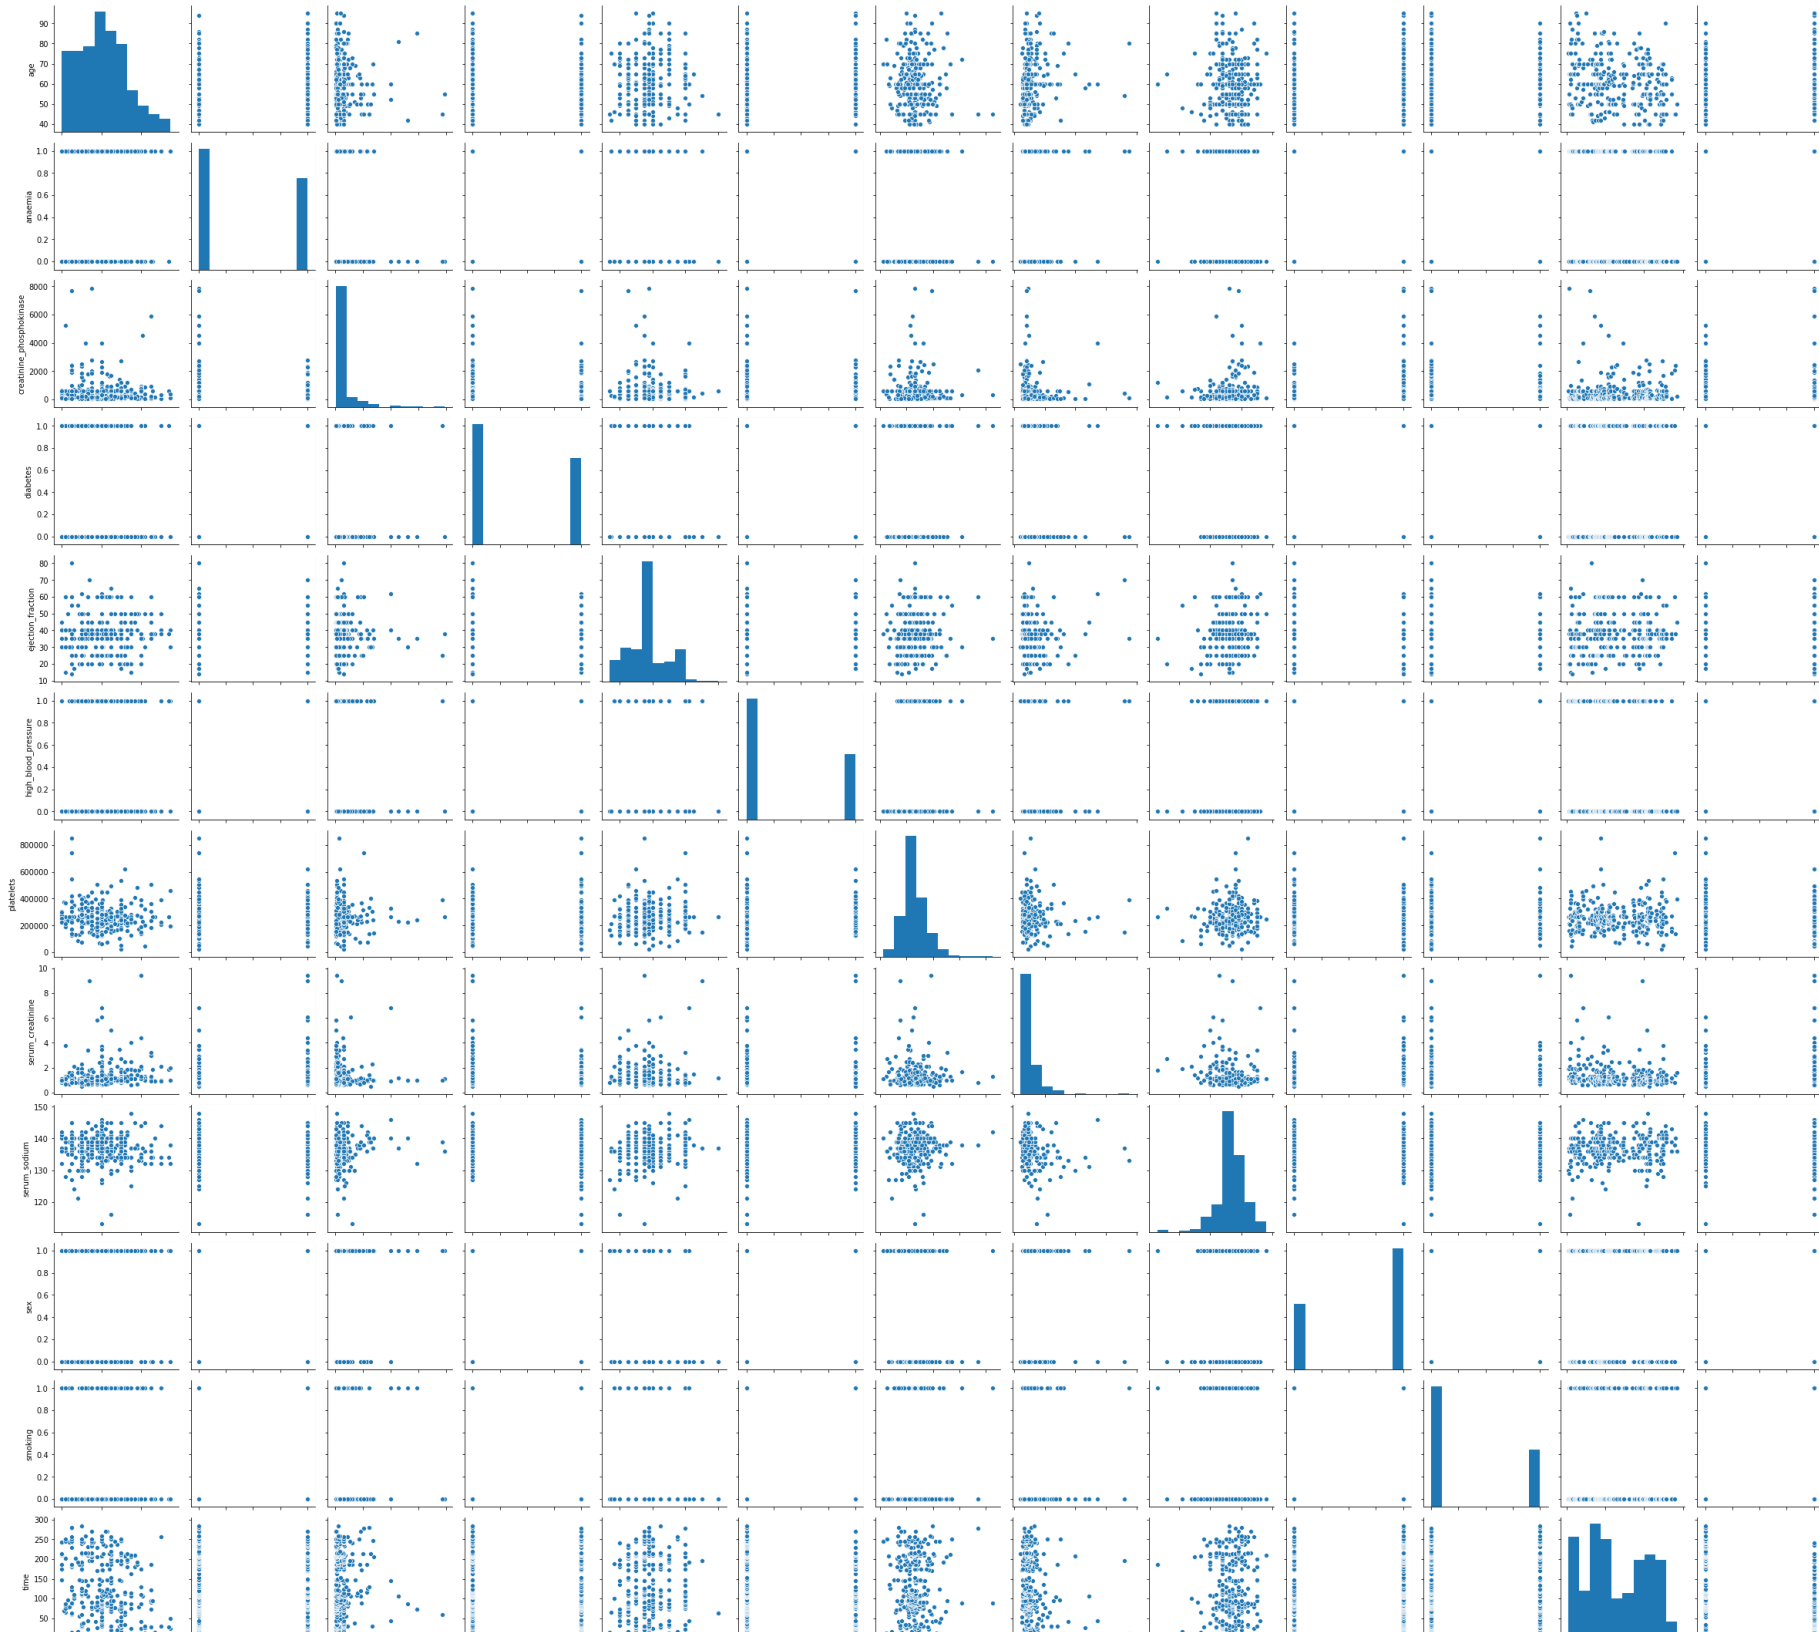
Out[7]: <bound method Index.unique of Index(['age', 'anaemia', 'creatinine\_phosphokinase', 'diabetes',  
 'ejection\_fraction', 'high\_blood\_pressure', 'platelets',  
 'serum\_creatinine', 'serum\_sodium', 'sex', 'smoking', 'time',  
 'DEATH\_EVENT'],  
 dtype='object')>

In [8]: *#comparing the correlation among features*


```
sns.pairplot(heart_failure)
```

Out[8]: <seaborn.axisgrid.PairGrid at 0x23313601e48>









In [9]: *#checking for correlation among our features using heatmap to visualize*

```
cmap = sns.diverging_palette(2, 165, s=80, l=55, n=9)
corrmat = heart_failure.corr()
plt.subplots(figsize=(20,20))
sns.heatmap(corrmat,cmap= cmap,annot=True, square=True)
```


Out[9]: <matplotlib.axes.\_subplots.AxesSubplot at 0x2331a4a9888>





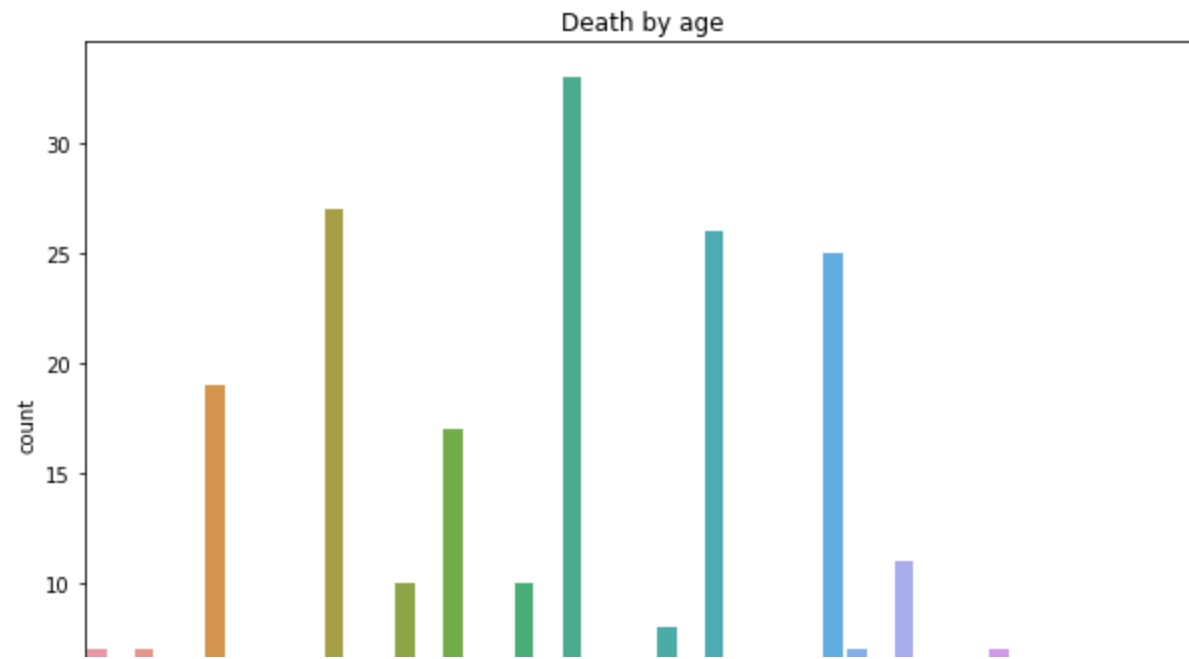
```
In [10]: #getting the total counts of our Label (Death event)
sns.countplot(x='DEATH_EVENT', data = heart_failure).set(title='Count of DEATH_EVENT')
```

```
Out[10]: [Text(0.5, 1.0, 'Count of DEATH_EVENT')]
```



```
In [11]: #to obtain the death event by age
```


```
plt.figure(figsize =(10,7))  
sns.countplot(x='age', data = heart_failure).set(title='Death by age')  
plt.xticks(rotation=90)  
plt.show()
```



- Most of the death occurred between the ages of 50 and 60

In [12]: *#box plot to determine possible outliers*

```
def plt_compare(continoues_features):  
    sns.boxplot(y=heart_failure[continoues_features],x=heart_failure.DEATH_EVENT)  
    plt.xticks(rotation=90)  
    plt.show()  
attributes=['age','creatinine_phosphokinase', 'ejection_fraction', 'platelets', 'serum_creatinine', 'serum_sodium', 'time_to_event']  
for attribute in attributes:  
    plt_compare(attribute)
```





For our continuous features,


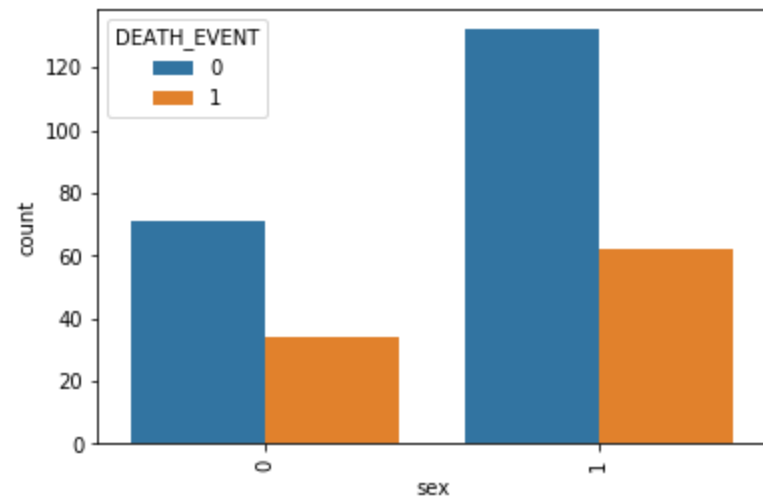

- we noticed all the plots had outliers however they will still be used in preprocessing

In [13]: *#To observe the effect of the categorical features on the death event caused by heart failure*

```
def plt_check(categorical_features):  
    sns.countplot(heart_failure[categorical_features], hue=heart_failure.DEATH_EVENT)  
    plt.xticks(rotation=90)  
    plt.show()  
attributes=['anaemia', 'diabetes', 'high_blood_pressure', 'sex', 'smoking']  
  
for attribute in attributes:  
    plt_check(attribute)
```







As seen, although, all categorical features contribute to the death event but appear not to have a strong significance in contributing to the target variable, we can decide to drop them as features and still be fine but we will work with them regardless.

In [14]:

```
# Split features into features and label (x,y)
X= heart_failure.drop([ 'DEATH_EVENT'], axis = 'columns')
X
```

Out[14]:

|     | age  | anaemia | creatinine_phosphokinase | diabetes | ejection_fraction | high_blood_pressure | platelets | serum_creatinine | serum_sodium | sex | si  |
|-----|------|---------|--------------------------|----------|-------------------|---------------------|-----------|------------------|--------------|-----|-----|
| 0   | 75.0 | 0       | 582                      | 0        | 20                | 1                   | 265000.00 | 1.9              | 130          | 1   |     |
| 1   | 55.0 | 0       | 7861                     | 0        | 38                | 0                   | 263358.03 | 1.1              | 136          | 1   |     |
| 2   | 65.0 | 0       | 146                      | 0        | 20                | 0                   | 162000.00 | 1.3              | 129          | 1   |     |
| 3   | 50.0 | 1       | 111                      | 0        | 20                | 0                   | 210000.00 | 1.9              | 137          | 1   |     |
| 4   | 65.0 | 1       | 160                      | 1        | 20                | 0                   | 327000.00 | 2.7              | 116          | 0   |     |
| ... | ...  | ...     | ...                      | ...      | ...               | ...                 | ...       | ...              | ...          | ... | ... |
| 294 | 62.0 | 0       | 61                       | 1        | 38                | 1                   | 155000.00 | 1.1              | 143          | 1   |     |
| 295 | 55.0 | 0       | 1820                     | 0        | 38                | 0                   | 270000.00 | 1.2              | 139          | 0   |     |
| 296 | 45.0 | 0       | 2060                     | 1        | 60                | 0                   | 742000.00 | 0.8              | 138          | 0   |     |
| 297 | 45.0 | 0       | 2413                     | 0        | 38                | 0                   | 140000.00 | 1.4              | 140          | 1   |     |
| 298 | 50.0 | 0       | 196                      | 0        | 45                | 0                   | 395000.00 | 1.6              | 136          | 1   |     |

299 rows × 12 columns



```
In [15]: y = heart_failure.DEATH_EVENT  
y
```

```
Out[15]: 0      1  
         1      1  
         2      1  
         3      1  
         4      1  
         ..  
        294     0  
        295     0  
        296     0  
        297     0  
        298     0  
        Name: DEATH_EVENT, Length: 299, dtype: int64
```

```
In [16]: from sklearn import preprocessing
```

```
In [17]: from sklearn.preprocessing import StandardScaler
```

```
In [18]: col_names = list(X.columns)
s_scaler = preprocessing.StandardScaler()
X_scaled= s_scaler.fit_transform(X)
X_scaled = pd.DataFrame(X_scaled, columns=col_names)
X_scaled.describe().T
```

Out[18]:

|                                 | count | mean          | std      | min       | 25%       | 50%       | 75%      | max      |
|---------------------------------|-------|---------------|----------|-----------|-----------|-----------|----------|----------|
| <b>age</b>                      | 299.0 | 5.265205e-16  | 1.001676 | -1.754448 | -0.828124 | -0.070223 | 0.771889 | 2.877170 |
| <b>anaemia</b>                  | 299.0 | 3.594301e-16  | 1.001676 | -0.871105 | -0.871105 | -0.871105 | 1.147968 | 1.147968 |
| <b>creatinine_phosphokinase</b> | 299.0 | 3.713120e-18  | 1.001676 | -0.576918 | -0.480393 | -0.342574 | 0.000166 | 7.514640 |
| <b>diabetes</b>                 | 299.0 | 1.113936e-16  | 1.001676 | -0.847579 | -0.847579 | -0.847579 | 1.179830 | 1.179830 |
| <b>ejection_fraction</b>        | 299.0 | 3.341808e-18  | 1.001676 | -2.038387 | -0.684180 | -0.007077 | 0.585389 | 3.547716 |
| <b>high_blood_pressure</b>      | 299.0 | -4.841909e-16 | 1.001676 | -0.735688 | -0.735688 | -0.735688 | 1.359272 | 1.359272 |
| <b>platelets</b>                | 299.0 | 1.009969e-16  | 1.001676 | -2.440155 | -0.520870 | -0.013908 | 0.411120 | 6.008180 |
| <b>serum_creatinine</b>         | 299.0 | -2.227872e-18 | 1.001676 | -0.865509 | -0.478205 | -0.284552 | 0.005926 | 7.752020 |
| <b>serum_sodium</b>             | 299.0 | -8.627435e-16 | 1.001676 | -5.363206 | -0.595996 | 0.085034  | 0.766064 | 2.582144 |
| <b>sex</b>                      | 299.0 | -5.940993e-18 | 1.001676 | -1.359272 | -1.359272 | 0.735688  | 0.735688 | 0.735688 |
| <b>smoking</b>                  | 299.0 | -3.861645e-17 | 1.001676 | -0.687682 | -0.687682 | -0.687682 | 1.454161 | 1.454161 |
| <b>time</b>                     | 299.0 | -1.069379e-16 | 1.001676 | -1.629502 | -0.739000 | -0.196954 | 0.938759 | 1.997038 |

```
In [19]: # import train, test, split library so as to split data into training and testing data
from sklearn.model_selection import train_test_split
```

```
In [20]: from sklearn.metrics import accuracy_score, confusion_matrix
```

```
In [21]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=101)
```

```
In [22]: from sklearn.svm import SVC
```

```
In [23]: from sklearn.ensemble import RandomForestClassifier
```

```
In [24]: from sklearn.linear_model import LogisticRegression
```

```
In [25]: # creating an array for our model of choice
models=[]
models.append(('LR',LogisticRegression(solver='liblinear',multi_class='ovr')))
models.append(('RFC',RandomForestClassifier()))
models.append(('SVM',SVC(gamma='auto')))
```

```
In [26]: from sklearn.model_selection import StratifiedKFold
```

```
In [27]: from sklearn.model_selection import cross_val_score
```

```
In [28]: results=[]
names=[]
for name,model in models:
    kfold=StratifiedKFold(n_splits=20)#random_state=1
    cv_results=cross_val_score(model, X_train, y_train, cv=kfold, scoring='accuracy')
    results.append(cv_results)
    names.append(name)
    print('%s: %f (%f)' %(name,cv_results.mean(),cv_results.std()))
```

```
LR: 0.812879 (0.084091)
RFC: 0.844318 (0.096794)
SVM: 0.660606 (0.035339)
```

- CV (cross validation): helps to evaluate and compare across different models
- Stratified Kfold: helps us to divide each of our classifier category in a uniform way
- kfold: we used 20 splits and took the average of each classifier alongside their standard deviation

```
In [29]: from sklearn.linear_model import LogisticRegression
```

```
In [30]: lr = LogisticRegression()
```

```
In [31]: lr.fit(X_train,y_train)
```

```
Out[31]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='auto', n_jobs=None, penalty='l2',
                             random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                             warm_start=False)
```

```
In [32]: #lets now predict the death event from heart failure
y_pred = lr.predict(X_test)
```

```
In [33]: lr.score(X_test,y_test)
```

```
Out[33]: 0.84
```

```
In [34]: from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report, accuracy_score, f1_score
```

```
In [35]: print(classification_report(y_test, y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.89      | 0.89   | 0.89     | 55      |
| 1            | 0.70      | 0.70   | 0.70     | 20      |
| accuracy     |           |        | 0.84     | 75      |
| macro avg    | 0.80      | 0.80   | 0.80     | 75      |
| weighted avg | 0.84      | 0.84   | 0.84     | 75      |

- accuracy is a score used to evaluate the models performance, the higher it is the better
- Recall measures the model ability to correctly predict the true positive values
- Precision is the ratio of true positive to the sum of both true and false positive
- f score combines precision and recall into one metric, its value should be closest to 1
- support is a number of actual occurrences of each class in the dataset


```
In [36]: from sklearn.metrics import accuracy_score, confusion_matrix
```

```
In [37]: # using cofusion metrics to check the behaviour of our Linear regression model
cm = confusion_matrix(y_test, y_pred)
cm
```

```
Out[37]: array([[49,  6],
               [ 6, 14]], dtype=int64)
```

```
In [38]: plt.figure(figsize=(10,7))
sns.heatmap(cm, annot = True)
plt.xlabel('predicted')
plt.ylabel('truth')
```

```
Out[38]: Text(69.0, 0.5, 'truth')
```



```
In [39]: svc_model= SVC(kernel = 'linear')
svc_model.fit(X_train, y_train)
```

```
Out[39]: SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='scale', kernel='linear',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```
In [40]: prediction= svc_model.predict(X_test)
```

```
In [41]: print(svc_model.score(X_train, y_train))
print(svc_model.score(X_test, y_test))
```

```
0.8125
0.8533333333333334
```

```
In [50]: print(classification_report(y_test, prediction))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.89      | 0.91   | 0.90     | 55      |
| 1            | 0.74      | 0.70   | 0.72     | 20      |
| accuracy     |           |        | 0.85     | 75      |
| macro avg    | 0.81      | 0.80   | 0.81     | 75      |
| weighted avg | 0.85      | 0.85   | 0.85     | 75      |


```
In [42]: # using confusion metrics to check the behaviour of svm model
cm2 = confusion_matrix(prediction, y_test)
cm2
```

```
Out[42]: array([[50,  6],
               [ 5, 14]], dtype=int64)
```



```
In [43]: plt.figure(figsize=(10,7))  
sns.heatmap(cm2, annot = True)  
plt.xlabel('predicted')  
plt.ylabel('truth')
```

```
Out[43]: Text(69.0, 0.5, 'truth')
```



```
In [44]: rf = RandomForestClassifier()  
rf.fit(X_train,y_train)  
pred_y=rf.predict(X_test)
```

```
In [45]: rf.score(X_test,y_test)
```

```
Out[45]: 0.8666666666666667
```

```
In [51]: print(classification_report(y_test, pred_y))
```


|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.94      | 0.87   | 0.91     | 55      |
| 1            | 0.71      | 0.85   | 0.77     | 20      |
| accuracy     |           |        | 0.87     | 75      |
| macro avg    | 0.82      | 0.86   | 0.84     | 75      |
| weighted avg | 0.88      | 0.87   | 0.87     | 75      |

```
In [46]: # using cofusion metrics to check the behaviour of our Random forest classifier model  
cm3 = confusion_matrix(pred_y, y_test)  
cm3
```

```
Out[46]: array([[48,  3],  
               [ 7, 17]], dtype=int64)
```

```
In [47]: plt.figure(figsize=(10,7))  
sns.heatmap(cm3, annot = True)  
plt.xlabel('predicted')  
plt.ylabel('truth')
```

```
Out[47]: Text(69.0, 0.5, 'truth')
```



3 models were used to predict the death event by heart failure. Confusion metrics was made for all the 3 models and the random forest classifier appear to give the best prediction and behaviour.